

Crafting In-context Examples according to LMs’ Parametric Knowledge

Yoonsang Lee^{◇,♡,*} Pranav Atreya^{♣,*} Xi Ye[◇] Eunsol Choi[◇]

[◇]The University of Texas at Austin, [♡]Seoul National University, [♣]UC Berkeley
lysianthus@snu.ac.kr, pranavatreya@berkeley.edu
xiye@cs.utexas.edu, eunsol@utexas.edu

Abstract

In-context learning can improve the performances of knowledge-rich tasks such as question answering. In such scenarios, in-context examples trigger a language model (LM) to surface information stored in its parametric knowledge. We study how to better construct in-context example sets, based on whether the model is aware of the in-context examples. We identify ‘known’ examples, where models can correctly answer from their parametric knowledge, and ‘unknown’ ones. Our experiments show that prompting with ‘unknown’ examples decreases the performance, potentially as it encourages hallucination rather than searching for its parametric knowledge. Constructing an in-context example set that presents both known and unknown information performs the best across diverse settings. We perform analysis on three multi-answer question answering datasets, which allows us to further study answer set ordering strategies based on the LM’s knowledge of each answer. Together, our study sheds light on how to best construct in-context example sets for knowledge-rich tasks.¹

1 Introduction

Large language models (LLMs) can perform competitively on knowledge-rich tasks such as question answering via in-context demonstrations (Brown et al., 2020). In such scenarios, in-context examples are used not only to teach the LLM the mapping from inputs to outputs, but also to invoke the LLM’s parametric knowledge (Liu et al., 2021; Agrawal et al., 2022). Given such role of in-context examples, we examine how the LLM’s parametric knowledge of in-context examples impact the effectiveness of in-context examples.

Let’s imagine a very challenging in-context example set, where LLMs cannot answer any of

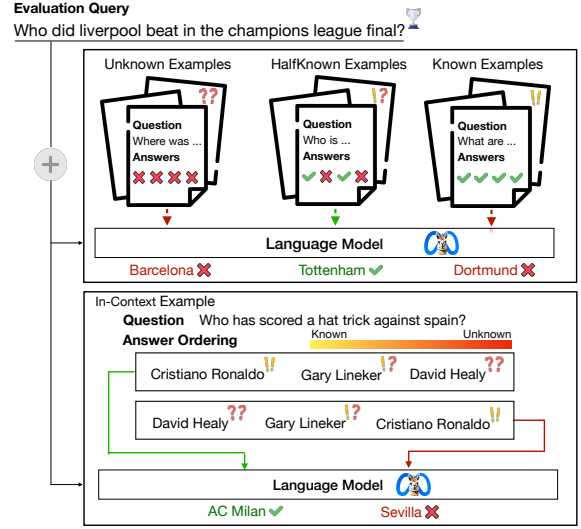


Figure 1: We study how an LM’s knowledge of in-context examples impacts its effectiveness. On the top box, we construct three sets of in-context examples, Unknown, HalfKnown, and Known, differing in their difficulty (Section 3). On the bottom box, we construct two in-context examples, which contain the same question and answer set, but the answers are sorted differently: one in increasing amount of parametric knowledge and one in reverse (Section 4, 5).

in-context examples from its parametric knowledge. For example, in-context examples can query knowledge about recent events that happened after pre-training (Lazaridou et al., 2021). These in-context examples may teach the model to generate plausible-looking responses, but also encourage hallucination as a result. On the other hand, if we only provide in-context examples where LLM can easily answer, would LLM learn to make an educated guess on more challenging evaluation examples?

We pose a suite of research questions connecting parametric knowledge of an LM on in-context examples and its impact on model predictions. Figure 1 provides our study overview. We mainly focus on multi-answer QA datasets (Min et al.,

^{*}Equal Contribution, work was done at UT Austin.

¹Our code is available at https://github.com/lilys012/known_examples.

2020; Malaviya et al., 2023; Amouyal et al., 2022) since they allow selecting and ordering the answers among multiple answers. Furthermore, we evaluate on a math QA dataset (Cobbe et al., 2021) and two NLI datasets (Dagan et al., 2005; Bowman et al., 2015), which require reasoning from LLMs.

We first compare providing ‘known’ or ‘unknown’ in-context examples (Section 3). We operationalize ‘known’ in-context examples as those LM can correctly predict with in-context learning. We do not observe a clear winner between two choices, with results varying depending on the dataset. Throughout all datasets, however, providing a set of in-context examples that have a mixture of known and unknown information leads to superior performance compared to a set consisting solely of known or unknown in-context examples.

Our next analysis focuses on the ordering of multi-answer set while fixing in-context example set (Section 4, 5). Compared to randomly ordering valid answers, semantically meaningful ordering brings substantial changes in model predictions. Even alphabetical ordering of answer set changes predicted answers substantially, prompting the model to generate 1.5 more answers on average than when shown a randomly sorted answer set. We further find that ordering the answer set of in-context examples in descending order of model knowledge often leads to performance gains. Together, our work suggests best practices for crafting in-context examples, with relation to their parametric knowledge, for knowledge-intensive tasks.

2 Experimental Settings

We first describe our evaluation setting which centers around multi-answer QA datasets.

2.1 Dataset

We evaluate on three multi-answer QA datasets: (1) AmbigQA (Min et al., 2020) contains a subset of questions from the Natural Questions (Kwiatkowski et al., 2019) dataset, namely those marked as ambiguous in the sense that depending on the interpretation, they can have multiple correct answers. (2) QAMPARI (Amouyal et al., 2022) consists of questions whose set of correct answers necessarily span multiple paragraphs in the document from which they were retrieved. The dataset was originally developed to evaluate retrieval methods, and we repurpose it to create a challenging closed-book QA setting.

(3) QUEST (Malaviya et al., 2023) dataset is constructed by formulating queries that define implicit set operations over Wikipedia entities. We report the dataset statistics in Appendix A.

2.2 Evaluation Metrics

Given a question q , the model will predict a set of answers $\hat{a} = \{a_1, a_2, \dots, a_m\}$, where each $a_i = (w_{i_1}, w_{i_2}, \dots, w_{i_{|a_i|}})$ is a sequence of tokens for a single answer. We denote $a^* = \{a_1^*, a_2^*, \dots, a_n^*\}$ as the ground truth answers to the same question.

We use standard token match metrics for evaluating answer accuracy, Exact Match (EM) and F1-score (Joshi et al., 2017). EM assigns a score of 1 if the predicted answer equals the ground truth answer, while F1-score is calculated over the tokens in the answer. We use metrics for multi-answers introduced in prior work (Min et al., 2020), which we describe below for completeness.

Answer-level Exact Match ($F1_{EM}$) As predicting the exact ground truth answer set correctly is very challenging, we report the F1-score of answer-level exact match, denoted as $F1_{EM}$. For an answer a and reference answers set S , we define a correctness score $c(f, a, S) = f(a, S)$ with respect to function f . We use $f(a, S) = \mathbb{1}(a \in S)$ here. Then, we calculate the F1-score over set-level precision and recall according to c .

$$P = \frac{\sum_{i=1}^m c(f, a_i, a^*)}{m}, R = \frac{\sum_{j=1}^n c(f, a_j^*, \hat{a})}{n}$$

$$F1_{EM} = \frac{2 \times P \times R}{P + R}$$

Answer-level F1 ($F1_{F1}$) The generated answer may be semantically equivalent to one of the ground truth answers, without being lexically equivalent (e.g., "Friends" and "The TV show Friends"). To account for such semantic equivalences, we use F1 score between the tokens of two answer strings instead of the exact match as a correctness score, $f(a, S) = \max_{a' \in S} (F1(a, a'))$. Then, we compute F1-score over set-level precision and recall as above.

Statistical Testing As our evaluation datasets are relatively small, we conduct paired bootstrap tests throughout most of our experiments, highlighting results that outperform baseline with p value of ≤ 0.05 .

	AmbigQA _{dev}		QAMPARI _{dev}		QUEST _{test}	
	Llama2	GPT-3.5	Llama2	GPT-3.5	Llama2	GPT-3.5
Random	18.0 / 28.9	20.0 / 31.6	10.3 / 20.8	15.0 / 28.5	3.4 / 11.0	6.0 / 16.6
Unknown	17.2* / 28.2*	20.3* / 33.1*	10.9* / 22.0*	14.8 / 27.9*	3.7* / 11.9*	5.7* / 15.8*
HalfKnown	18.5* / 29.5*	21.6* / 33.2*	11.3* / 22.6	15.5* / 28.2*	4.0* / 11.9*	6.3* / 17.4*
Known	18.3* / 29.0*	21.3* / 33.1*	9.8 / 19.7	15.3 / 29.2*	3.9* / 12.0*	5.4* / 15.8

Table 1: Results comparing known example and unknown example. We present $F1_{EM}$ and then $F1_{F1}$ in each cell. Using half-known example outperforms other settings. We put * on scores that are significantly different from that of Random in-context examples set, and bold the highest performing set for each metric.

2.3 Base Models

Language Model We evaluate on Llama2 (Touvron et al., 2023) (13B) language model mainly, and additionally OPT (Zhang et al., 2022) (13B) and GPT-3.5-turbo models to evaluate generalization.

In-Context Example Retriever Prior work (Rubin et al., 2021) has established that using semantically similar in-context examples improves the performance of in-context learning significantly. Throughout our study, we often retrieve top 5 most similar in-context examples from the entire training set for each dataset to form the prompt. We place in-context examples in increasing order of similarity, such that the most similar example will be presented the closest to the evaluation question. We measure example similarities by encoding each question with a SimCSE model (Gao et al., 2021) and computing their dot product.

3 Known Examples vs. Unknown Examples

Prior work has studied a few characteristics of successful in-context example set, such as label distribution in the in-context example set (Min et al., 2022). We evaluate in-context examples with respect to model’s parametric knowledge, whether a ‘known’ or ‘unknown’ in-context example is better. We operationalize ‘known’ ones as the ones where LLMs can get the answers correctly from its own parametric knowledge, and ‘unknown’ ones as those that the model did not answer it correctly.

3.1 In-Context Example Set Study

We create four sets of in-context examples, differing in its difficulty for a given LM.

- **UNKNOWN:** examples for which the LM possesses no knowledge of the answers. Operationally, these are examples when LM is

prompted with five most similar examples, LM will predict zero answers correctly (i.e. zero $F1_{EM}$ score).

- **RANDOM:** randomly sampled examples. Since the LM possesses no knowledge to majority of the examples, these exhibit 0.18 $F1_{EM}$ score on average.
- **HALFKNOWN:** examples for which the LM possesses roughly half knowledge of the answers (i.e. 0.5 $F1_{EM}$ score).
- **KNOWN:** examples for which the LM possesses full knowledge of the answers (i.e. 1.0 $F1_{EM}$ score).

As prior work (Rubin et al., 2021) has established that the similarity of in-context example to the query correlates strongly with the model’s performance, we control for this confounding factor. We compute the average similarity for each in-context example candidate to other in-context example candidates in the candidate set (training set). Then, we choose a fixed number of in-context examples whose average similarity value is close to the median value.² From this candidate set, we sample five examples for each condition and use them as fixed in-context examples across all questions in the evaluation dataset. To further reduce randomness, we sample multiple sets of five example set for each condition and report the average performance (by default, four sets are sampled and two sets are sampled for HALFKNOWN and KNOWN set in QUEST because of lack of examples with sufficient model knowledge).

We present the performance of each in-context example set for three datasets with Llama2 and GPT-3.5 in Table 1. We observe the HALF-

²We choose 999 examples for AmbigQA and QAMPARI, and 499 for QUEST (as QUEST only has 1251 training examples), half from below median, half from above median. For QUEST, we could not find enough examples with where model score full $F1_{EM}$ score, so we selected highest scoring examples. The mid-range is (0.245, 0.264), (0.294, 0.296), (0.326, 0.373) for AmbigQA, QAMPARI, and QUEST.




Figure 2: Results of single answer study on Llama2 model. Only an answer at the x -th quantile of perplexities in decreasing order is presented in each in-context example. As the model gets exposed to more known answers, the performance tend to increase.

KNOWN in-context example set achieves strong performance consistently on both LMs. Since half-known examples contain both answers that the model knows and doesn’t know, we hypothesize this may successfully prompt LMs to leverage parametric knowledge and to make educated guesses.

3.2 Single Answer Study

In this study, we further control for variability in the questions used in in-context examples. We fix the in-context example set and manipulate the multi-answer set, such that we provide only one answer from multi answer set for each in-context example. For example, if a question in in-context example is “Who was the president of U.S.?”, we can either provide a famous president or a lesser-known president as an answer. Both are ‘correct’ answers, but which answer would lead to better model performance?

For each question in our evaluation set, we retrieve the top five most similar examples from the training set as in-context examples. We measure the perplexity of each answer to approximate how well LM ‘knows’ the answer. For each example, a pair of question q and gold answer set $\{a_1^*, a_2^*, \dots, a_n^*\}$, we form a prefix p by prepending the top five most similar examples to the query q .³ Then, we compute the length normalized perplexity of each answer a_i^* and prefix p as follows:

$$PP(a_i^*|p) = \prod_{j=1}^{|a_i^*|} P(w_{ij}|p, w_{i1}, \dots, w_{i(j-1)})^{-\frac{1}{|a_i^*|}}$$

³We present an example prefix in Appendix F.

We order the gold answer set in descending order of perplexity, and select an answer at the x -th quantile. This way, an answer at the 100% quantile represents the most ‘known’ answer, as its perplexity is the lowest among the gold answers.

Figure 2 presents the $F1_{EM}$ score among various x -th quantile. We observe a clear trend across all three datasets, that using a ‘known’ answer leads LM to generate more accurate answer. The overall F1 score is low, due to low recall, as these in-context examples are incomplete.⁴ Nonetheless, this experiment affirms that crafting in-context examples by considering the model’s parametric knowledge can impact the final performance.

3.3 Extension to Other Tasks

In this section, we explore the generalizability of our findings from Section 3.1 to other tasks, single-answer QA and Natural Language Inference (NLI).

Datasets For single-answer QA task, we select the GSM8K (Cobbe et al., 2021) dataset, which is commonly used to assess the reasoning capabilities of LLMs. For NLI task, we choose RTE (Dagan et al., 2005) and SNLI (Bowman et al., 2015) datasets, two standard NLI benchmarks. For all datasets, we use their standard train, development, and test splits.

We first evaluate each example in training set, to identify whether LM’s parametric knowledge is sufficient to evaluate individual example correctly. We classify each example as correct, wrong, or invalid, where invalid indicates that the model did not produce an answer. For GSM8K, we use a fixed 8-shot example set taken from Wei et al. (2022b). For NLI task, for each training example, we retrieve the top five most similar example from the training set (excluding itself) to form a 5-shot example set. We define the similarity as the dot product of two SimCSE embedding. After this process, each training example is labeled as correct, wrong or invalid. We throw away invalid examples and sample from correct and wrong set to form in-context example set, of varying degrees of difficulty.

In-Context Example Set Unlike multi-answer QA, where examples can be partially correct, in these tasks, the examples are evaluated as either correct or incorrect. Therefore, we construct HALF-KNOWN set by mixing easier and harder in-context examples as follows:

⁴Since we provide only one answer for all in-context examples, LM predicts a single answer for each question.

	Unknown	Random	HalfKnown	Known
GSM8K _{test}	33.1	34.8	36.4	32.0
RTE _{dev}	72.2	79.1	79.8	79.1
SNLI _{test}	62.7	69.3	71.0	68.6

Table 2: Performance(accuracy) on GSM8K, RTE, and SNLI datasets. Accuracy is expressed as the percentage of correct answers over the entire test dataset.

- **UNKNOWN** set includes randomly selected six examples that model answered incorrectly.
- **RANDOM** set includes randomly selected six examples from entire training dataset.
- **HALFKNOWN** set includes three correct and three wrong examples.
- **KNOWN** set includes randomly selected six examples that model answered correctly.

Result We select six examples four times and report the averaged accuracy with Llama2 model in Table 2. On all three datasets, HALFKNOWN set achieves the highest accuracy, repeating the trend from multi-answer QA datasets.

4 Ordering Answers Based on LM’s Knowledge

Prior work suggests that the ordering of in-context examples significantly impacts the performance, with more relevant examples being most beneficial when placed last (Zhao et al., 2021). Yet, no prior work has studied how the ordering of answer set in in-context examples affects model generation and task performances. We investigate this here. Following our previous study, our focus is on **parametric** knowledge of LMs being prompted. Specifically, we question whether placing answers based on how well the model knows about answers improves the performance.

We present strategies to order the answer set of each example, a pair of question q and its gold answer set $a^* = \{a_1^*, a_2^*, \dots, a_n^*\}$, which will be used as an in-context example.⁵ We present two baselines and two methods (PERPLEXITY, GREEDY) for ordering the gold answer set of each in-context example based on model’s parametric knowledge.

Baselines The RANDOM baseline randomly orders answers and ALPHABET orders answers alphabetically. While alphabetical ordering is not

⁵As reordering process is computationally expensive, proportional to the number of answers, we only consider examples that have less than 20 answers. This results in exclusion of 1 example in AmbigQA, 8094 examples in QAMPARI, and none in QUEST.

Input: LM \mathcal{M} , Prefix p , Gold answer set $a^* = \{a_1^*, \dots, a_n^*\}$, where each gold answer is a token sequence (i.e., $a_i^* = (w_{i_1}, \dots, w_{i_{|a_i^*|}})$)

Output: Ordered answer indices of the gold answer set

```

1:  $I_1 \leftarrow \{w_{1_1}, \dots, w_{n_1}\}$ 
2:  $u \leftarrow 1$ 
3: while  $I_1 \neq \emptyset$  do
4:    $t \leftarrow 0$ 
5:   repeat
6:      $t \leftarrow t + 1$ 
7:      $o_t \leftarrow \operatorname{argmax}_{w \in I_t} P_{\mathcal{M}}(w|p)$ 
8:      $p \leftarrow [p; o_t]$ 
9:      $I_{t+1} \leftarrow \{w_{i_{t+1}} | w_{i_t} == o_t\}$ 
10:    until  $\exists a_{k_u}^* == (o_1, \dots, o_t)$  {this assigns  $k_u$  the index of completed answer}
11:    $I_1 \leftarrow I_1 \setminus \{w_{k_{u-1}}\}$ 
12:    $u \leftarrow u + 1$ 
13: return  $\{k_1, \dots, k_n\}$ 

```

Figure 3: Algorithm for constrained decoding for GREEDY ordering.

relevant to model’s parametric knowledge of the answer, prior work (Madaan et al., 2022) has shown that consistent ordering of labels can improve the performance of fine-tuned LLM’s predictions.

Knowledge-Aware Ordering We decide ordering based on the **perplexity** of individual answer given the prefix, or by performing **greedy** constrained decoding, given the prefix, on a smaller, restricted vocabulary set. We use the same prefix as in Section 3.2, a concatenation of five in-context examples. Each ordering strategy will yield two orderings of answers, which either sorts the answers in the descending order of model’s parametric knowledge or ascending order (denoted as REVERSE).

- **PERPLEXITY:** We compute the length normalized perplexity of each answer a_i^* and prefix p as used in Section 3.2. Then, we sort the answers in ascending order of these perplexities, resulting in ‘known’ answers placed earlier.
- **GREEDY:** We arrange the gold answers by performing a beam search decoding in a greedy manner, constrained to permissible tokens. There will be two loops, outer loop for selecting the first token of the generated answer, and inner loop for completing the chosen first token.

Figure 3 presents the pseudocode, which we explain below. Let’s denote a_i^* as a sequence of tokens $(w_{i_1}, w_{i_2}, \dots, w_{i_{|a_i^*|}})$ for the i -th answer. At each decoding step t , a set of permissible tokens I_t is constructed. Initially, $I_1 =$

		S				
		GREEDY	REVERSE GREEDY	PERPLEXITY	REVERSE PERPLEXITY	ALPHABET
\mathcal{D}_e	AmbigQA _{dev}	71.7 / 66.0	39.2 / 37.2	69.5 / 65.8	38.1 / 34.2	87.4 / 55.6
	QAMPARI _{dev}	69.6 / 60.0	42.2 / 41.0	58.1 / 54.1	46.3 / 45.9	95.0 / 58.9
	QAMPARI _{test}	70.0 / 65.7	43.0 / 41.7	58.8 / 55.8	45.0 / 44.2	94.9 / 58.1
	QUEST _{dev}	78.4 / 63.9	47.2 / 45.8	57.1 / 51.5	49.3 / 48.5	95.7 / 52.1
	QUEST _{test}	81.0 / 63.3	45.7 / 45.3	57.6 / 52.5	48.8 / 47.5	95.6 / 50.8
Average		74.1	43.5	60.2	45.5	93.7

Table 3: Percentage of generated answer ordering matching in-context examples answer ordering, where we use Llama2 for \mathcal{M} . In each cell, we present the percentage from using corresponding answer ordering strategy first ($\phi(S, \mathcal{D}_t^S, \mathcal{D}_e, \mathcal{M})$) and the percentage for randomly ordering answers for control ($\phi(S, \mathcal{D}_t^{\text{random}}, \mathcal{D}_e, \mathcal{M})$).

$\{w_{1_1}, w_{2_1}, \dots, w_{n_1}\}$, a set of the first token for each potential answer. We choose a token from this set that has the highest likelihood given the prompt, i.e. $o_1 = \operatorname{argmax}_{w \in I_1} P(w|p)$. Then, we update the prefix $p \leftarrow [p; o_1]$. This initiates the inner loop, setting $I_2 = \{w_{i_2} | w_{i_1} = o_1\}$ as a set of second token of answers who starts with the selected first token. This continues until one of the answers a_{k_1} is fully generated. Afterwards, we come back to the outer loop, and the initial set of permissible tokens is set to be $I_1 = \{w_{1_1}, w_{2_1}, \dots, w_{n_1}\} \setminus \{w_{t_1}\}$ excluding $a_{k_1}^*$ which has been already generated. This process continues until all answers has been generated, with a time complexity of $O(n|a_i^*|)$.

5 Results for Answer Ordering Strategies

Having introduced strategies for ordering answers for in-context examples, we study how this impacts the generation of answers with Llama2 and OPT. We first evaluate whether the generated answers mimic the ordering of answers in in-context examples. Then, we evaluate whether the ordering impacts the size and the accuracy of predicted answer set. We also report whether two model’s parametric knowledges are in sync, meaning, if one model knows about one fact, does the other model likely to know the same fact? We overall observe such patterns, particularly for QUEST dataset.

5.1 Does the predicted answer set follow the ordering of in-context answer set?

Throughout in-context learning, the model is expected to learn the pattern shown in the demonstrations (Min et al., 2022). We assess the generated answers to observe if the model has followed the particular ordering shown in in-context examples.

Metric We introduce a metric $\phi(S, \mathcal{D}_t^S, \mathcal{D}_e, \mathcal{M})$. This measures how much LM \mathcal{M} follows the an-

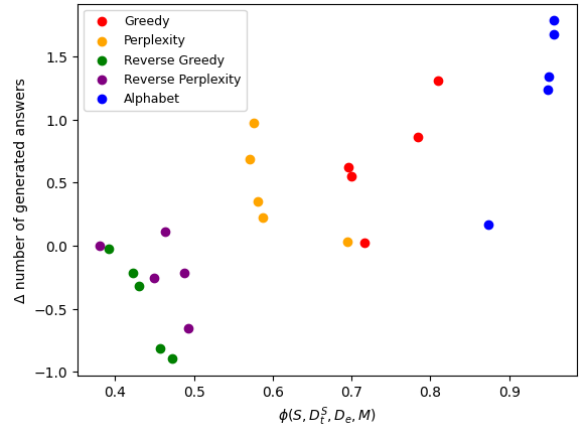


Figure 4: $\phi(S, \mathcal{D}_t^S, \mathcal{D}_e, \mathcal{M})$ vs. the number of generated answers across three datasets, where we use Llama2 for \mathcal{M} . Instead of the raw number of answer set, we report the size difference compared to the answer set generated from random ordering. As ϕ increases, which signifies how faithfully LM follows the ordering strategy in in-context examples, the model generates more answers.

swer ordering strategy S on evaluation dataset \mathcal{D}_e when using in-context examples from training dataset \mathcal{D}_t whose answered are ordered according to S^t .⁶ When S matches S^t , this metric will measure how much predicted outputs mimic the answer ordering strategy of in-context examples.

Let’s denote $\hat{a}_i = \{a_{i_1}, a_{i_2}, \dots, a_{i_m}\}$ be the list of predicted m answers for the i -th example of an evaluation dataset \mathcal{D}_e , following its generation order from model \mathcal{M} . We reorder the predicted answers from \hat{a}_i with respect to S and denote $f(a_{ij})$ to be the index of a_{ij} in the newly ordered set.

For each consecutive answer pair in \hat{a}_i , we evaluate whether their order is preserved after reordering. Then we count the number of consecutive answer pairs that have preserved the ordering, which is $P_i = \sum_{j=1}^{m-1} \mathbb{1}(f(a_{ij}) < f(a_{i(j+1)}))$. Similarly,

⁶We assume retrieving five most similar in-context examples for each evaluation example throughout this study.

<i>AmbigQA</i>	P_{EM}	R_{EM}	$F1_{EM}$	$F1_{F1}$
RANDOM	27.1	17.9	20.0	31.3
GREEDY	27.2	18.5	20.5	31.7
PERPLEXITY	27.4	18.4	20.5	31.8
REVERSE GREEDY	27.1	17.8	20.1	31.5
REVERSE PERPLEXITY	27.3	17.9	20.2	31.8
ALPHABET	26.7	18.2	20.3	31.2

<i>QAMPARI</i>	P_{EM}	R_{EM}	$F1_{EM}$	$F1_{F1}$
RANDOM	26.3 / 25.2	11.7 / 10.9	13.8 / 12.9	25.3 / 22.4
GREEDY	26.4 / 25.7	12.2 / 11.9*	14.2 / 14.0*	25.6 / 22.6
PERPLEXITY	26.7 / 26.4*	12.4* / 11.6*	14.6* / 13.9*	25.8 / 22.9
REVERSE GREEDY	26.5 / 25.8	11.6 / 10.1*	13.9 / 12.4	25.1 / 21.8
REVERSE PERPLEXITY	27.0 / 26.7*	11.7 / 11.0	14.0 / 13.3	25.2 / 22.5
ALPHABET	24.5* / 23.5*	12.7* / 11.8*	14.3 / 13.6	24.7 / 22.6

<i>QUEST</i>	P_{EM}	R_{EM}	$F1_{EM}$	$F1_{F1}$
RANDOM	23.9 / 24.8	17.9 / 19.7	18.3 / 19.9	27.2 / 27.8
GREEDY	23.8 / 24.8	19.6* / 20.8*	19.5* / 20.6*	28.6* / 28.4*
PERPLEXITY	24.3 / 24.8	19.3* / 20.8*	19.4 / 20.6*	28.0 / 28.4*
REVERSE GREEDY	22.9 / 24.5	17.0 / 18.4*	17.4 / 18.8*	26.3 / 26.5*
REVERSE PERPLEXITY	23.7 / 24.5	17.3 / 19.4	17.7 / 19.4	26.4 / 27.1*
ALPHABET	20.5* / 23.8*	17.6 / 20.4*	17.0 / 20.0	25.0* / 27.0*

Table 4: QA performance for answer ordering strategies on Llama2 (13B) model. P_{EM} and R_{EM} are precision and recall for calculating $F1_{EM}$. We present development set performance and then test set performance in each cell. Blue color indicates improved performance compared to Random and red indicates the opposite. We put * on scores that are significantly different from that of Random ordering.

$N_i = \sum_{j=1}^{m-1} \mathbb{1}(f(a_{ij}) > f(a_{i(j+1)}))$ represents the number of pairs that violate the ordering. Then, we compute micro average over \mathcal{D}_e .

$$\phi(S, \mathcal{D}_t^{St}, \mathcal{D}_e, \mathcal{M}) = \frac{100 \cdot \sum_{i \in \mathcal{D}_e} P_i}{\sum_{i \in \mathcal{D}_e} (P_i + N_i)}$$

Results Table 3 presents the results for Llama2 model, and we provide the results for OPT model in Table 8 in the appendix. For each $\phi(S, \mathcal{D}_t^S, \mathcal{D}_e, \mathcal{M})$, we also report $\phi(S, \mathcal{D}_t^{S_{random}}, \mathcal{D}_e, \mathcal{M})$ as a control. We found that in every cell (except for one cell in Table 8), the first number is higher than the second number, suggesting that the model follows the answer ordering pattern presented in the in-context examples. We found this is particularly true for ALPHABET ordering, which is probably the easiest pattern to learn.

We further observe that the model is decoding answers such that it will present **confident** answer first (following the orders of GREEDY and PERPLEXITY), even when answers in in-context example is randomly ordered. Even after introducing consistent ordering (presenting less confident answer first), the model shows propensity to present confident answer first (values for REVERSE GREEDY and REVERSE PERPLEXITY are below chance (50) consistently).

5.2 Does ordering impact the number of generated answers?

Unlike in simpler QA tasks where there is exactly one gold answer, models have to decide how many answers to generate. Would consistent ordering of answers allow the model to generate more answers?

We report the number of generated answers for each ordering strategy for Llama2 model in Figure 4. We find that generation order impacts the number of generated answer, with ALPHABET ordering substantially increasing the number of generated answers the most. The results further suggest that an ordering pattern that is easier for the model to learn can prompt LM to generate more answers. We report the results for OPT model in Figure 8 which shows the same trends.

5.3 Does the ordering impact the QA performance?

Lastly, we examine the end task (QA) performance of different answer ordering strategies. Table 4 presents the results on Llama2 model. Overall, we see that answer ordering does not bring large impact in final performance, but notice consistent patterns. Presenting more confident answers first (GREEDY and PERPLEXITY) yielded better results than their REVERSE counterparts. GREEDY and




Figure 5: Plots of log answer perplexities from Llama2-13b (x-axis) and OPT-13b (y-axis). Horizontal and vertical lines indicate the mean value of log perplexities with respect to each LM. In all datasets, Llama2 outperforms OPT in its parametric knowledge, and the answers mostly report higher perplexity with OPT compared to Llama2.

PERPLEXITY show gains mostly in recall, leading to increase in both $F1_{EM}$ and $F1_{F1}$. Arbitrary, yet consistent ordering such as ALPHABET does not improve model performance, sometimes rather leading to lower performance. The trend holds for AmbigQA though not statistically significant. This might be caused by smaller average answer set size compared to that of other datasets (2-3 vs. 10+ answers). We suggest ordering ‘known’ answer first in in-context examples to improve model performance.

For OPT model, we observe GREEDY and PERPLEXITY show improved performance through gains in recall for QUEST dataset but the results are mostly random on other datasets (Table 9 in the appendix). We plot the perplexity of individual answer in train examples with respect to two models in Figure 5. Overall, we find that Llama2 contains more factual knowledge than OPT, resulting in higher end task performance. Two models exhibit similar knowledge for QUEST as they strongly correlate, however OPT shows a wider range of perplexities on other datasets, especially for answers that have low perplexity on Llama2. We hypothesize carefully ordering between answers will bring significant changes in end task performance **only** when model exhibits sufficient parametric knowledge of subset of answers. When the model is not familiar enough with the gold answers in in-context examples, knowledge-aware answer ordering might have limited effectiveness.

5.4 Transfer to other base LMs

So far we have measured the parametric knowledge on an language model and then use the same model for in-context prompting. In this section, we experiment using in-context example set constructed with parametric knowledge of one language model

(Llama2), see how it impacts the generation of another language model (GPT-3.5). While different LMs have different pre-training data, the relative parametric knowledge might be similar for different LMs (e.g., famous entity to one LM remains famous for another LM). This also allows us to experiment with propriety black-box LM API easily, whose prediction probability is not always available. We observe similar patterns as in the original experiments (Table 5), but the effect size is much smaller and not significant, potentially because of the difference in parametric knowledge between two models.

5.5 Random In-Context Examples

Prior works have highlighted the importance of relevant in-context examples, such as those based on similarity (Liu et al., 2021) and diversity (Levy et al., 2022). Yet, many studies do not do example specific retrieval and use random examples for its simplicity. Throughout our experiments in Section 5, we retrieved similar in-context examples for each evaluation example. How would our results hold if we use randomly select in-context examples?

First, with randomly retrieved in-context examples, models still learn to follow the answer ordering strategy shown in in-context examples but substantially less than when using similar in-context examples (Table 10 in the appendix). Second, we find that the number of generated answer is affected similarly, with using ALPHABET ordering leads to the highest number of generated answers. However, we see invariant performances on end tasks (Table 11 in the appendix). Carefully constructing relevant in-context examples is more meaningful than doing it for random in-context examples. This suggests that if you do not have large enough training examples to recover semantically relevant in-context

<i>AmbigQA</i>	P_{EM}	R_{EM}	$F1_{EM}$	$F1_{F1}$
RANDOM	28.2	22.1	23.1	35.7
PERPLEXITY	28.8	23.1*	23.9	36.5
REVERSE PERPLEXITY	29.0	22.3	23.5	35.3
ALPHABET	28.4	22.5	23.5	35.8
<i>QAMPARI</i>	P_{EM}	R_{EM}	$F1_{EM}$	$F1_{F1}$
RANDOM	23.4 / 23.2	18.7 / 18.5	18.4 / 18.4	30.1 / 28.4
PERPLEXITY	23.9 / 22.9	19.5 / 19.1	18.9 / 18.5	30.4 / 29.1
REVERSE PERPLEXITY	23.2 / 23.1	18.2 / 18.5	18.2 / 18.3	30.2 / 28.5
ALPHABET	23.4 / 23.0	17.3* / 17.8	17.8 / 18.0	29.0* / 27.5
<i>QUEST</i>	P_{EM}	R_{EM}	$F1_{EM}$	$F1_{F1}$
RANDOM	15.0 / 16.4	16.7 / 17.6	14.8 / 15.8	25.5 / 26.4
PERPLEXITY	16.6 / 17.0	17.7 / 18.6*	15.9 / 16.5*	26.8 / 26.8
REVERSE PERPLEXITY	16.2 / 16.5	17.5 / 17.8	15.5 / 15.9	26.6 / 26.4
ALPHABET	15.5 / 17.0	16.2 / 17.6	14.9 / 16.2	24.9 / 25.5*

Table 5: QA performance for answer ordering strategies with GPT-3.5 model. We order the answer set with respect to parametric knowledge of Llama2 and evaluate its transfer to GPT-3.5 model.

examples, careful construction of prompt might not yield changes in end task performance.

6 Related Work

Multi Label Ordering While not studied extensively under the in-context-learning setting, a recent work (Madaan et al., 2022) studies set generation problem from an encoder-decoder model, showing that imposing informative ordering over the label space improves model performance.

Analysis on In-context Learning Many prior works investigate factors that determine the performance of in-context learning (Brown et al., 2020), such as the composition of the pre-training dataset (Xie et al., 2022), size of language model (Wei et al., 2022a), number of pre-training tokens (Touvron et al., 2023), and specific fine-tuning strategy employed (Wei et al., 2021). More closely related to ours, one line of work particularly focuses on factors related to the in-context examples, including the choice of verbalizer and templates (Min et al., 2022), order of examples (Lu et al., 2022; Pezeshkpour and Hruschka, 2023), and the choice of in-context examples (Liu et al., 2021; Rubin et al., 2021; Agrawal et al., 2022; Ye et al., 2023). While past work is mainly centered around classification tasks, our work studies the task of multi-answer QA, with a focus on how LM’s parametric knowledge on in-context examples impact the performance. In particular, our findings suggests that answers with lower perplexity lead to more accurate answer, which is congruent with recent work that shows using lower perplexity prompts im-

proves model perplexity in general (Ye and Durrett, 2023; Iter et al., 2023; Gonen et al., 2022).

Multi-answer QA Real-world questions could naturally have multiple answers when a question is ambiguous (Min et al., 2020; Stelmakh et al., 2022), when a question is evaluated under different temporal or geographical contexts (Zhang and Choi, 2021), or when a question expects a set of answers (Amouyal et al., 2022; Malaviya et al., 2023). While most prior work tackles multi-answer QA in the open-book setting by retrieving from external corpus (Shao and Huang, 2022; Sun et al., 2023), we study the problem in the close-book setting, which prompts LLMs to generate the answers based on their parametric knowledge.

7 Conclusion

We present comprehensive studies on knowledge-aware prompt design for multi-answer QA tasks. Our findings underscore the benefits of having in-context examples that the language model is familiar with. First, the HALFKNOWN set aids the model in effectively accessing its parametric knowledge. Second, employing knowledge-aware ordering of presenting answers in descending order of the model’s knowledge enhances the overall process of answer generation.

Limitations

Our study mainly focuses on multi-answer QA datasets, while we have demonstrated the potential for generalization to single-answer QA and NLI tasks. The analysis can be extended to a wide range of tasks that requires various types of reasoning abilities. Also, we find that the end task performance gets less impacted when random in-context examples are used. Further studies could explore diverse in-context example retrieval methods as well as cover multiple languages.

Acknowledgment

We thank the members of UT Austin NLP community for valuable feedback, especially Thom Lake for suggesting an ablation study and Greg Durrett for providing feedback on the draft. This work is partially funded by a gift from Home Depot. This work was supported by Korea Institute for Advancement of Technology (KIAT) grant funded by the Korea Government (Ministry of Education) (P0025681-G02P22450002201-10054408, Semiconductor-Specialized University).

References

- Sweta Agrawal, Chunting Zhou, Mike Lewis, Luke Zettlemoyer, and Marjan Ghazvininejad. 2022. In-context examples selection for machine translation. In *Annual Meeting of the Association for Computational Linguistics*.
- Samuel Joseph Amouyal, Ohad Rubin, Ori Yoran, Tomer Wolfson, Jonathan Herzig, and Jonathan Berant. 2022. Qampari: An open-domain question answering benchmark for questions with many answers from multiple paragraphs. *ArXiv, abs/2205.12665*.
- Samuel R. Bowman, Gabor Angeli, Christopher Potts, and Christopher D. Manning. 2015. A large annotated corpus for learning natural language inference. In *Conference on Empirical Methods in Natural Language Processing*.
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. 2020. Language models are few-shot learners. *Advances in neural information processing systems*, 33:1877–1901.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. 2021. Training verifiers to solve math word problems. *ArXiv, abs/2110.14168*.
- Ido Dagan, Oren Glickman, and Bernardo Magnini. 2005. The pascal recognising textual entailment challenge. In *Machine Learning Challenges Workshop*.
- Tianyu Gao, Xingcheng Yao, and Danqi Chen. 2021. Simcse: Simple contrastive learning of sentence embeddings. *ArXiv, abs/2104.08821*.
- Hila Gonen, Srinu Iyer, Terra Blevins, Noah A. Smith, and Luke Zettlemoyer. 2022. Demystifying prompts in language models via perplexity estimation. *ArXiv, abs/2212.04037*.
- Dan Iter, Reid Pryzant, Ruochen Xu, Shuo Wang, Yang Liu, Yichong Xu, and Chenguang Zhu. 2023. In-context demonstration selection with cross entropy difference. *ArXiv, abs/2305.14726*.
- Mandar Joshi, Eunsol Choi, Daniel S. Weld, and Luke Zettlemoyer. 2017. Triviaqa: A large scale distantly supervised challenge dataset for reading comprehension. *ArXiv, abs/1705.03551*.
- Tom Kwiatkowski, Jennimaria Palomaki, Olivia Redfield, Michael Collins, Ankur Parikh, Chris Alberti, Danielle Epstein, Illia Polosukhin, Jacob Devlin, Kenton Lee, et al. 2019. Natural questions: a benchmark for question answering research. *Transactions of the Association for Computational Linguistics*, 7:453–466.
- Angeliki Lazaridou, Adhiguna Kuncoro, Elena Gribovskaya, Devang Agrawal, Adam Liska, Tayfun Terzi, Mai Gimenez, Cyprien de Masson d’Autume, Tomás Kociský, Sebastian Ruder, Dani Yogatama, Kris Cao, Susannah Young, and Phil Blunsom. 2021. Mind the gap: Assessing temporal generalization in neural language models. In *Neural Information Processing Systems*.
- Itay Levy, Ben Bogin, and Jonathan Berant. 2022. Diverse demonstrations improve in-context compositional generalization. *ArXiv, abs/2212.06800*.
- Jiachang Liu, Dinghan Shen, Yizhe Zhang, Bill Dolan, Lawrence Carin, and Weizhu Chen. 2021. What makes good in-context examples for gpt-3? *arXiv preprint arXiv:2101.06804*.
- Yao Lu, Max Bartolo, Alastair Moore, Sebastian Riedel, and Pontus Stenetorp. 2022. Fantastically ordered prompts and where to find them: Overcoming few-shot prompt order sensitivity. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*.
- Aman Madaan, Dheeraj Rajagopal, Niket Tandon, Yiming Yang, and Antoine Bosselut. 2022. Conditional set generation using seq2seq models. In *Conference on Empirical Methods in Natural Language Processing*.
- Chaitanya Malaviya, Peter Shaw, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2023. Quest: A retrieval dataset of entity-seeking queries with implicit set operations. *arXiv preprint arXiv:2305.11694*.

- Sewon Min, Xinxu Lyu, Ari Holtzman, Mikel Artetxe, Mike Lewis, Hannaneh Hajishirzi, and Luke Zettlemoyer. 2022. Rethinking the role of demonstrations: What makes in-context learning work? *arXiv preprint arXiv:2202.12837*.
- Sewon Min, Julian Michael, Hannaneh Hajishirzi, and Luke Zettlemoyer. 2020. Ambigqa: Answering ambiguous open-domain questions. *arXiv preprint arXiv:2004.10645*.
- Pouya Pezeshkpour and Estevam Hruschka. 2023. Large language models sensitivity to the order of options in multiple-choice questions. *ArXiv*, abs/2308.11483.
- Ohad Rubin, Jonathan Herzig, and Jonathan Berant. 2021. Learning to retrieve prompts for in-context learning. *arXiv preprint arXiv:2112.08633*.
- Zhihong Shao and Minlie Huang. 2022. Answering open-domain multi-answer questions via a recall-then-verify framework. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1825–1838, Dublin, Ireland. Association for Computational Linguistics.
- Ivan Stelmakh, Yi Luan, Bhuwan Dhingra, and Ming-Wei Chang. 2022. ASQA: Factoid questions meet long-form answers. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 8273–8288, Abu Dhabi, United Arab Emirates. Association for Computational Linguistics.
- Weiwei Sun, Hengyi Cai, Hongshen Chen, Pengjie Ren, Zhumin Chen, Maarten de Rijke, and Zhaochun Ren. 2023. Answering ambiguous questions via iterative prompting. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 7669–7683, Toronto, Canada. Association for Computational Linguistics.
- Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, et al. 2023. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*.
- Jason Wei, Maarten Bosma, Vincent Y Zhao, Kelvin Guu, Adams Wei Yu, Brian Lester, Nan Du, Andrew M Dai, and Quoc V Le. 2021. Finetuned language models are zero-shot learners. *arXiv preprint arXiv:2109.01652*.
- Jason Wei, Yi Tay, Rishi Bommasani, Colin Raffel, Barret Zoph, Sebastian Borgeaud, Dani Yogatama, Maarten Bosma, Denny Zhou, Donald Metzler, et al. 2022a. Emergent abilities of large language models. *arXiv preprint arXiv:2206.07682*.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, et al. 2022b. Chain-of-thought prompting elicits reasoning in large language models. *Advances in Neural Information Processing Systems*, 35:24824–24837.
- Sang Michael Xie, Aditi Raghunathan, Percy Liang, and Tengyu Ma. 2022. An explanation of in-context learning as implicit bayesian inference. In *International Conference on Learning Representations*.
- Jiacheng Ye, Zhiyong Wu, Jiangtao Feng, Tao Yu, and Lingpeng Kong. 2023. Compositional exemplars for in-context learning. *arXiv preprint arXiv:2302.05698*.
- Xi Ye and Greg Durrett. 2023. Explanation selection using unlabeled data for in-context learning. In *Proceedings of EMNLP*.
- Michael Zhang and Eunsol Choi. 2021. SituatedQA: Incorporating extra-linguistic contexts into QA. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 7371–7387, Online and Punta Cana, Dominican Republic. Association for Computational Linguistics.
- Susan Zhang, Stephen Roller, Naman Goyal, Mikel Artetxe, Moya Chen, Shuohui Chen, Christopher Dewan, Mona T. Diab, Xian Li, Xi Victoria Lin, Todor Mihaylov, Myle Ott, Sam Shleifer, Kurt Shuster, Daniel Simig, Punit Singh Koura, Anjali Sridhar, Tianlu Wang, and Luke Zettlemoyer. 2022. Opt: Open pre-trained transformer language models. *ArXiv*, abs/2205.01068.
- Tony Zhao, Eric Wallace, Shi Feng, Dan Klein, and Sameer Singh. 2021. Calibrate before use: Improving few-shot performance of language models. In *International Conference on Machine Learning*.

A Dataset Statistics

We report the dataset statistics in Table 6.

B Similarity of In-Context Examples

We calculate the similarity score of two in-context examples using SimCSE embeddings of each query. Figure 6 illustrates the similarity distributions across three datasets.




Figure 6: Similarity distributions among in-context example candidates. The x-axis denotes embedding similarity (with SimCSE (Gao et al., 2021) encoder) and the y-axis indicates the percentage of each bin. The median value for each dataset is 0.254, 0.295, 0.350.

C Experimental Details

C.1 Resources

All experiments are conducted on NVIDIA A40 GPU. A single evaluation for AmbigQA and QUEST (development split) took around 20 minutes. QAMPARI (development and test split) took around 1 hours. QUEST (test split) took around 2 hours, due to its largest size.

C.2 Statistical Testing

We conduct paired bootstrap tests with 10000 bootstrap samples throughout our experiments (Section 2.2). Since we have multiple (two or four) in-context example sets for experiments in Section 3, we randomly sample one in-context example set of each class (UNKNOWN, HALFKNOWN, KNOWN, and RANDOM) and conduct testing.

D In-Context Example Set Study

In Table 7, we present the results from Section 3.1 for QAMPARI_{test} and QUEST_{dev} on Llama2.

E Answer Ordering Strategies

E.1 Single Answer Study

We examine the effectiveness of answer ordering strategies discussed at Section 4. We provide only one answer at the forefront of each ordered answers in in-context examples. Since an answer from GREEDY and PERPLEXITY is ‘known’ to the model, they may serve as an upper bound of ‘known’ answer, while REVERSE GREEDY and REVERSE PERPLEXITY may serve as a lower bound. RANDOM exists somewhere between these. The disparities among these are clear, as shown in Figure 7. The results suggest that the model is able to differentiate ordering strategies.




Figure 7: Answer-level Exact Match ($F1_{EM}$) score for demonstrating only one frontmost answer of an ordering methodology on Llama2 model.

E.2 Results on OPT 13B model

We present the results of experiments in Section 5 with OPT 13B model. With respect to following the ordering strategy of in-context examples (Section 5.1, 5.2), we find that the results hold for OPT LLM model as well (Table 8). However, the end task performance results are somewhat mixed (Table 9, Figure 8). We observe consistent results of end task performance on QUEST dataset but the results are mostly random on AmbigQA and QAMPARI dataset.

E.3 Results on GPT-3.5 model

GPT-3.5-turbo model tends to generate lengthy and chatty outputs such as “There is not enough information given to answer this question”. Therefore we add a short instruction as following: “Follow the answers pattern”.

	AmbigQA		QAMPARI			QUEST		
	Train	Dev.	Train	Dev.	Test	Train	Dev.	Test
# Examples	4,615	1,048	50,372	1,000	1,000	1,251	316	1,669
Avg. # of answers	2.8	3.1	14.0	13.2	13.1	10.9	10.7	10.7
Query length	46.9	46.7	67.8	57.7	55.8	54.0	52.2	53.3
Answer length	15.9	14.5	14.4	17.3	16.6	17.2	16.7	17.0
Answer sequence length	45.2	45.4	200.9	228.5	217.6	187.0	179.0	182.4
# Unique answers	10,684	2,999	455,469	12,462	12,464	10,160	3,050	12,367

Table 6: Dataset statistics. Lengths of query, answer, and answer sequence are measured by the length of each string. # Unique answers counts unique answers within each split. Duplicated questions are removed from training sets.

	QAMPARI _{test}		QUEST _{dev}	
	$F1_{EM}$	$F1_{F1}$	$F1_{EM}$	$F1_{F1}$
Random	10.0	19.3	4.0	12.1
Unknown	10.6	20.2*	4.4*	13.2*
HalfKnown	11.2*	20.9*	4.9*	13.1*
Known	9.9	18.6	4.3*	12.8*

Table 7: Results comparing known example and unknown example. We put * on scores that are significantly different from that of Random in-context examples set, and bold the highest performing set for each metric.




Figure 8: $\phi(S, \mathcal{D}_t^S, \mathcal{D}_e, \mathcal{M})$ vs. the number of generated answers across three datasets, where we use OPT (13B) model for \mathcal{M} .

F Prompts

Throughout Table 12 to Table 17, we present the prompts used in our experiments.

		<i>S</i>				
		GREEDY	REVERSE GREEDY	PERPLEXITY	REVERSE PERPLEXITY	ALPHABET
\mathcal{D}_e	AmbigQA _{dev}	60.3 / 58.3	43.7 / 42.2	68.8 / 58.1	49.5 / 41.9	75.5 / 50.5
	QAMPARI _{dev}	62.8 / 52.1	39.0 / 39.6	60.0 / 55.1	52.1 / 44.9	87.8 / 52.0
	QAMPARI _{test}	63.1 / 52.4	39.7 / 39.1	61.8 / 56.7	52.1 / 39.1	85.4 / 47.3
	QUEST _{dev}	70.5 / 49.1	44.0 / 42.5	60.0 / 57.1	53.1 / 42.9	91.1 / 67.6
	QUEST _{test}	75.3 / 57.5	46.3 / 45.5	60.1 / 54.0	50.6 / 46.0	92.3 / 51.6
Average		66.4	42.5	62.1	51.5	86.4

Table 8: Percentage of generated answer ordering matching in-context examples answer ordering, where we use OPT (13B) model for \mathcal{M} . The table is formatted the same as Table 3.

<i>AmbigQA</i>	P_{EM}	R_{EM}	$F1_{EM}$	$F1_{F1}$
RANDOM	13.1	10.3	10.7	19.4
GREEDY	13.1	10.3	10.7	19.5
PERPLEXITY	12.9	10.0	10.5	19.2
REVERSE GREEDY	12.9	9.9	10.5	19.1
REVERSE PERPLEXITY	13.2	10.7	11.0	19.3
ALPHABET	13.5	10.6	11.0	19.3
<i>QAMPARI</i>	P_{EM}	R_{EM}	$F1_{EM}$	$F1_{F1}$
RANDOM	14.2 / 15.5	7.5 / 7.2	8.1 / 8.2	18.6 / 17.1
GREEDY	14.0 / 14.9	7.5 / 7.6	7.9 / 8.4	18.6 / 17.8
PERPLEXITY	14.7 / 15.6	7.8 / 7.7	8.3 / 8.5	19.0 / 17.6
REVERSE GREEDY	14.5 / 15.4	6.9* / 6.7	7.6 / 7.9	18.0* / 16.7
REVERSE PERPLEXITY	15.6* / 15.9	7.6 / 7.2	8.4 / 8.3	18.8 / 16.9
ALPHABET	14.4 / 15.0	8.1* / 7.9*	8.5 / 8.9*	18.7 / 17.4
<i>QUEST</i>	P_{EM}	R_{EM}	$F1_{EM}$	$F1_{F1}$
RANDOM	14.6 / 18.4	11.6 / 16.1	12.0 / 15.6	21.3 / 23.8
GREEDY	15.7 / 18.6	16.6* / 18.0*	14.9* / 17.0*	23.7* / 25.2*
PERPLEXITY	16.1 / 18.3	14.8* / 17.0*	13.9* / 16.2	22.6 / 24.5
REVERSE GREEDY	14.5 / 17.4*	10.7 / 13.8*	10.2 / 13.8*	19.6 / 22.1*
REVERSE PERPLEXITY	15.0 / 17.9	14.3* / 15.4	13.2 / 15.1	22.4 / 23.5
ALPHABET	16.3* / 17.6*	15.9* / 17.3*	14.7* / 16.3*	23.0* / 24.1*

Table 9: QA performance for answer ordering strategies with OPT (13B) model. The table is formatted the same as Table 4.

		<i>S</i>				
		GREEDY	REVERSE GREEDY	PERPLEXITY	REVERSE PERPLEXITY	ALPHABET
\mathcal{D}_e	AmbigQA _{dev}	69.6 / 68.9	33.7 / 32.8	70.2 / 70.5	68.9 / 29.5	83.6 / 62.5
	QAMPARI _{dev}	63.2 / 59.8	40.7 / 40.3	57.0 / 57.3	57.3 / 42.7	92.6 / 65.9
	QAMPARI _{test}	61.2 / 61.4	43.5 / 43.2	57.5 / 56.4	57.5 / 43.6	92.7 / 60.7
	QUEST _{dev}	55.4 / 52.6	39.3 / 40.2	59.1 / 57.4	57.1 / 42.6	88.5 / 59.7
	QUEST _{test}	56.8 / 54.0	38.9 / 40.1	56.9 / 56.4	56.4 / 43.6	86.7 / 60.9
Average		61.2	39.2	60.1	59.4	88.8

Table 10: Percentage of generated answer ordering matching in-context examples answer ordering, where we employ random in-context examples instead of most similar examples. The table is formatted the same as Table 3.

	<i>AmbigQA_{dev}</i>			<i>QAMPARI_{dev}</i>			<i>QAMPARI_{test}</i>		
	<i>F1_{EM}</i>	<i>F1_{F1}</i>	# ans	<i>F1_{EM}</i>	<i>F1_{F1}</i>	# ans	<i>F1_{EM}</i>	<i>F1_{F1}</i>	# ans
RANDOM	17.8	28.7	2.07	9.8	20.2	3.77	10.0	19.1	3.74
GREEDY	17.4	27.8	2.12	9.6	19.9	4.42	9.3	17.7	4.43
PERPLEXITY	17.9	28.3	2.11	9.7	20.0	3.99	9.7	18.6	4.03
REVERSE GREEDY	17.6	28.3	2.08	9.8	20.4	3.82	9.6	18.5	3.61
REVERSE PERPLEXITY	17.9	28.4	2.11	9.3	19.7	3.83	9.6	18.4	3.81
ALPHABET	17.9	28.5	2.22	9.8	19.8	5.48	9.6	17.5	5.41

	<i>QUEST_{dev}</i>			<i>QUEST_{test}</i>		
	<i>F1_{EM}</i>	<i>F1_{F1}</i>	# ans	<i>F1_{EM}</i>	<i>F1_{F1}</i>	# ans
RANDOM	4.4	12.9	3.42	3.5	11.2	3.41
GREEDY	4.7	12.5	4.51	3.4	10.9	4.49
PERPLEXITY	4.7	13.0	3.60	3.4	11.1	3.62
REVERSE GREEDY	4.0	12.5	3.51	3.3	11.1	3.11
REVERSE PERPLEXITY	4.6	12.6	3.09	3.6	11.4	3.28
ALPHABET	4.5	11.2	5.84	3.0	9.4	5.99

Table 11: QA performance for answer ordering strategies with random in-context examples. We bold the highest performing set for each metric.

Question	Who is the current chairman of african union commission?
Gold Answers	Jean Ping, Moussa Faki, Nkosazana Clarice Dlamini-Zuma
Prompt	Question: Who is the chairman of the federal reserve?\nAnswers: Alan Greenspan Ben Bernanke Janet Yellen\n\nQuestion: Who is the president of south africa now?\nAnswers: Thabo Mvuyelwa Mbeki Kgalema Petrus Motlanthe JZ\n\nQuestion: Who is the present chairperson of national human rights commission in india?\nAnswers: Justice K. G. Balakrishnan H. L. Dattu Cyriac Joseph\n\nQuestion: Who appoints the chairman of the finance commission?\nAnswers: the President Pranab Mukherjee Ram Nath Kovind Pratibha Devisingh Patil\n\nQuestion: Who is the chairman of national commission for woman of india?\nAnswers: Lalitha Kumaramangalam Mamta Sharma Girija Vyas\n\nQuestion: Who is the current chairman of african union commission?\nAnswers:
Output	Jean Ping Nkosazana Dlamini-Zuma Moussa Faki Mahamat\n

Table 12: Prompt example of AmbigQA

Question	What movies did Scott Z. Burns screenwrite?
Gold Answers	Contagion, No Time to Die, Pu-239, Side Effects, The Bourne Ultimatum, The Informant!, The Laundromat, The Mercy, The Report
Prompt	Question: Which film has Edward Burns as a member of its cast and had Edward Burns as screenwriter?\nAnswers: Ash Wednesday Purple Violets She's the One Sidewalks of New York The Brothers McMullen The Groomsman Newlyweds Looking for Kitty No Looking Back\n\nQuestion: Scott Ziehl was a director for what genre of film?\nAnswers: action film crime thriller horror film monster film drama\n\nQuestion: What are the publication dates of film that had Scott Z. Burns as screenwriter?\nAnswers: 2007 2006 2009 2013 2014 2018 2019 2020\n\nQuestion: Who are the cast members of film that had Scott Z. Burns as screenwriter?\nAnswers: Christoph Waltz Patton Oswalt Marsha Stephanie Blake Jocko Sims David Costabile Dan Fredenburgh Ted Levine Ken Stott Steven Berkoff\n\nQuestion: What are the genres of film that had Scott Z. Burns as screenwriter?\nAnswers: action film thriller film comedy film spy film drama\n\nQuestion: What movies did Scott Z. Burns screenwrite?\nAnswers:
Output	Contagion Side Effects The Bourne Ultimatum The Bourne Supremacy The Bourne Identity\n

Table 13: Prompt example of QAMPARI

Question	Mary Stewart novels
Gold Answers	A Walk in Wolf Wood, Airs Above the Ground, My Brother Michael, The Crystal Cave, The Hollow Hills, The Ivy Tree, The Last Enchantment, The Prince and the Pilgrim, The Wicked Day, The Wind Off the Small Isles, This Rough Magic, Thornyhold, Touch Not the Cat, Wildfire at Midnight
Prompt	Question: 1740s fantasy novels, 1747 books, or novels by Denis Diderot\nAnswers: Niels Klim's Underground Travels The Sofa: A Moral Tale Jacques the Fatalist Rameau's Nephew La Religieuse The Art of Cookery Made Plain and Easy Letters from a Peruvian Woman Madame de La Carli\nQuestion: Sarah Weeks novels\nAnswers: So B. It Jumping the Scratch\nQuestion: 1858 British novels, 1872 fantasy novels, or Novels by George MacDonald\nAnswers: At the Back of the North Wind Alec Forbes of Howglen The Princess and the Goblin David Elginbrod The Day Boy and the Night Girl Ranald Bannerman's Boyhood The Princess and Curdie Through the Looking-Glass Phantastes My Lady Ludlow Sir Gibbie Mary Marston Doctor Thorne Lilith Poor Dionis\nQuestion: Mark Clapham novels\nAnswers: The Taking of Planet 5 Twilight of the Gods The Book of the War Beige Planet Mars Hope\nQuestion: Marlon James Novels\nAnswers: A Brief History of Seven Killings John Crow's Devil The Book of Night Women\nQuestion: Mary Stewart novels\nAnswers:
Output	The Crystal Cave The Hollow Hills The Last Enchantment The Wicked Day\n

Table 14: Prompt example of QUEST

Question	Every day, Wendi feeds each of her chickens three cups of mixed chicken feed, containing seeds, mealworms and vegetables to help keep them healthy. She gives the chickens their feed in three separate meals. In the morning, she gives her flock of chickens 15 cups of feed. In the afternoon, she gives her chickens another 25 cups of feed. How many cups of feed does she need to give her chickens in the final meal of the day if the size of Wendi's flock is 20 chickens?
Logical reasoning and Answer	If each chicken eats 3 cups of feed per day, then for 20 chickens they would need $3 \times 20 = 60$ cups of feed per day.\nIf she feeds the flock 15 cups of feed in the morning, and 25 cups in the afternoon, then the final meal would require $60 - 15 - 25 = 20$ cups of chicken feed.\n#### 20
Prompt	Question: Mabel lives 4500 steps directly east of Lake High school. Helen lives $\frac{3}{4}$ the number of steps that Mabel lives, directly west of the school. What's the total number of steps Mabel will walk to visit Helen so that they can do their assignments together?\nAnswer: Helen lives $\frac{3}{4} \times 4500 = 3375$ steps directly west of Lake High. To reach Helen, Mabel would have to walk to $4500 + 3375 = 7875$ steps. #### 7875\nQuestion: Mark is 7 years older than Amy, who is 15. How old will Mark be in 5 years?\nAnswer: Mark is 15 years + 7 years = $15 + 7 = 22$ years old. In 5 years, he will be 22 years + 5 years = $22 + 5 = 27$ years old. #### 27\nQuestion: Steve has 2 boxes of pencils with 12 pencils in each box. He gave Matt 3 more pencils than he gave to Lauren. If Steve gave 6 pencils to Lauren, how many pencils does he have left?\nAnswer: Steve started with $2 \times 12 = 24$ pencils. He gave Matt $6 + 3 = 9$ pencils. After giving away the pencils, Steve will have $24 - 9 = 15$ pencils left. #### 15\nQuestion: Mandy researched 42 med schools. She applied to $\frac{1}{3}$ of the schools she researched and got into half of the schools where she applied. How many schools was Mandy accepted to?\nAnswer: First find the number of schools Mandy applied to: $42 \text{ med schools} \div 3 = 14$ med schools Then divide that number by 2 to find the number of schools where she was accepted: $14 \text{ med schools} \div 2 = 7$ med schools #### 7\nQuestion: Rachel is stuffing envelopes. She has eight hours to complete the task, and there are 1,500 envelopes. In the first hour, Rachel stuffs 135 envelopes. The second hour she stuffs 141 envelopes. How many envelopes will Rachel need to stuff per hour to finish the job?\nAnswer: Rachel has $1500 - 135 - 141 = 1224$ envelopes remaining to stuff. Rachel has 8 hours - 2 hours = 6 hours left to finish the task. Rachel needs to stuff $1224 \text{ envelopes} \div 6 \text{ hours} = 204$ envelopes per hour. #### 204\nQuestion: Samson is going to another town which is 140 km away. He will use his car that uses ten liters of gasoline for a distance of 70 km. How many liters of gasoline will Samson need for a one-way trip?\nAnswer: Samson will need $140 \text{ km} \div 70 \text{ km} = 2$ ten liters of gasoline for a one-way trip to a town. Therefore, he will need a total of $2 \times 10 \text{ liters} = 20$ liters of gasoline. #### 20\nQuestion: Every day, Wendi feeds each of her chickens three cups of mixed chicken feed, containing seeds, mealworms and vegetables to help keep them healthy. She gives the chickens their feed in three separate meals. In the morning, she gives her flock of chickens 15 cups of feed. In the afternoon, she gives her chickens another 25 cups of feed. How many cups of feed does she need to give her chickens in the final meal of the day if the size of Wendi's flock is 20 chickens?\nAnswer:
Output	Wendi gives her chickens 15 cups of feed in the morning and 25 cups of feed in the afternoon. She needs to give her chickens another 20 cups of feed in the final meal of the day. #### 20

Table 15: Prompt example of GSM8K

Premise	Yet, we now are discovering that antibiotics are losing their effectiveness against illness. Disease-causing bacteria are mutating faster than we can come up with new antibiotics to fight the new variations.
Hypothesis	Bacteria is winning the war against antibiotics.
Answer	entail
Prompt	Please identify whether the premise entails the hypothesis. The answer should be exact 'entail' or 'not entail'. <pre> \npremise: A brain-dead teenager was removed from life-support equipment after giving birth to a premature daughter. \nhypothesis: Pregnant women are allowed to die after they deliver their babies. \nanswer: not entail. \npremise: Kuwait becomes the fourth Gulf country where women as well as men can vote in elections. Saudi Arabia holds strictly limited, male-only elections. \nhypothesis: Women can vote in elections in Saudi Arabia. \nanswer: not entail. \npremise: The first Windows DNA lab outside Microsoft was established in suburban Philadelphia in June 1998. \nhypothesis: Microsoft was established in June 1998. \nanswer: not entail. \npremise: Scientists at the Genome Institute of Singapore (GIS) have discovered the complete genetic sequence of a coronavirus isolated from a Singapore patient with SARS. \nhypothesis: Singapore scientists reveal that SARS virus has undergone genetic changes. \nanswer: not entail. \npremise: Two portraits, one by Reynolds, the other by Gainsborough, since 1990 among the most looked-after works in the Art Loss Register database, were recovered this week when they were brought into Sotheby's for evaluation. Also a lithograph by Norwegian artist Edvard Munch, stolen from an Oslo art gallery in April, has been recovered, news reports said today. \nhypothesis: These two portraits and a lithography by Edvard Munch were recovered using the Art Loss Register database. \nanswer: not entail. \npremise: The deal marks the second foray into Lithuania by Philip Morris which in April beat British American Tobacco to acquire the Klaipeda Tobacco company for Dollars 40m - the biggest western investment in the Baltic states so far. \nhypothesis: The acquisition of the Klaipeda Tobacco company is the biggest investment in the Baltic states. \nanswer: not entail. \npremise: Yet, we now are discovering that antibiotics are losing their effectiveness against illness. Disease-causing bacteria are mutating faster than we can come up with new antibiotics to fight the new variations. \nhypothesis: Bacteria is winning the war against antibiotics. \nanswer: </pre>
Output	entail.

Table 16: Prompt example of RTE

Premise	This church choir sings to the masses as they sing joyous songs from the book at a church.
Hypothesis	A choir singing at a baseball game.
Answer	no
Prompt	Please identify whether the premise entails the hypothesis. The answer should be exact 'yes', 'no' or 'neutral'. <pre> \npremise: A man skis past another man displaying paintings in the snow. \nhypothesis: A man skis past his brother. \nanswer: neutral. \npremise: A little boy holding a yellow, plastic shovel, crouches in the sand. \nhypothesis: A little girl is holding a shovel and crouches in the sand. \nanswer: no. \npremise: A little girl wearing a yellow coat, striped scarf, pink rain boots, and carrying a white purse is walking with her Golden Retriever down an icy and puddled road. \nhypothesis: A little girl, wearing a yellow coat, is walking her dog down a road. \nanswer: yes. \npremise: Some African American young adults are playing volleyball. \nhypothesis: People playing basketball. \nanswer: no. \npremise: A brown and black dog is laying on a shaggy rug. \nhypothesis: A white dog is lying on a wooden floor. \nanswer: no. \npremise: Two black and white dogs in a field of flowers and grass. \nhypothesis: Some animals are outdoors. \nanswer: yes. \npremise: This church choir sings to the masses as they sing joyous songs from the book at a church. \nhypothesis: A choir singing at a baseball game. \nanswer: </pre>
Output	no.

Table 17: Prompt example of SNLI